

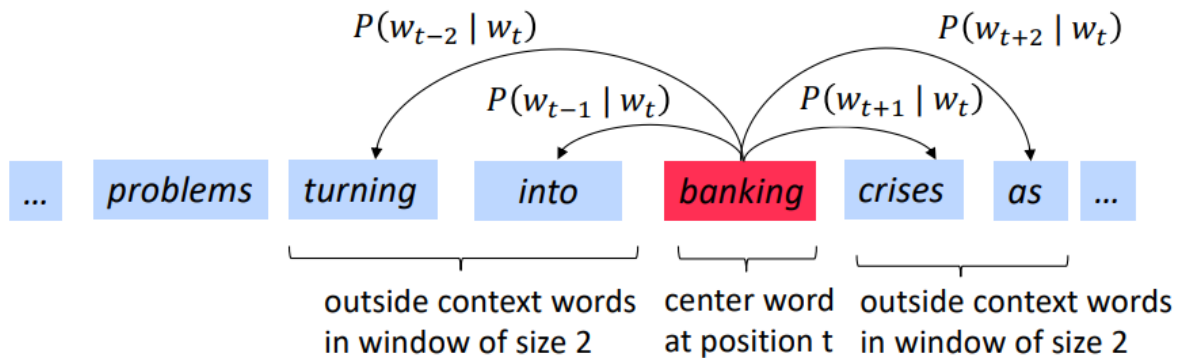
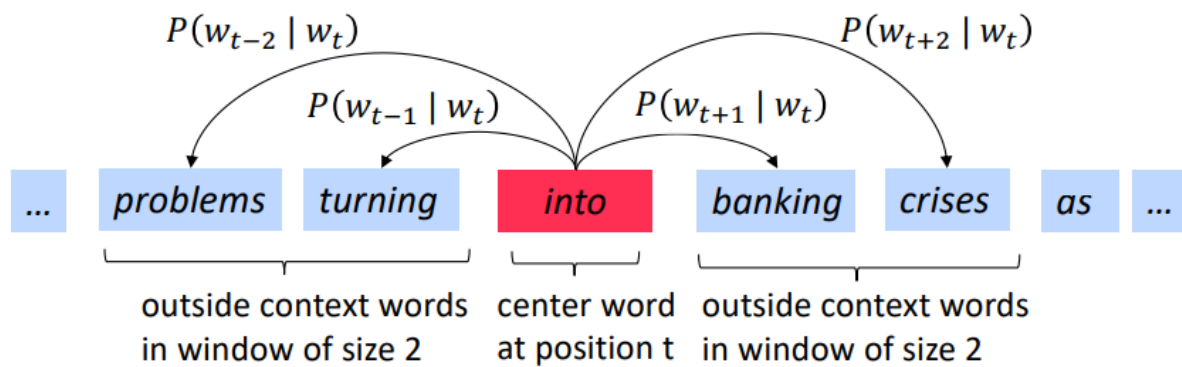
FastText

Word2Vec Idea:

Word2vec (Mikolov et al. 2013) is a framework for learning word vectors Idea:

- We have a large corpus of text
- Every word in a fixed vocabulary is represented by a vector
- Go through each position t in the text, which has a center word c and context (“outside”) words o
- Use the similarity of the word vectors for c and o to calculate the probability of o given c (or vice versa)
- Keep adjusting the word vectors to maximize this probability

Example windows and process for computing: $P(w_{t+j} | w_t)$



Objective Function:

For each position $t = 1, \dots, T$, predict context words within a window of fixed size m , given center word w_j .

$$\text{Likelihood} = L(\theta) = \prod_{t=1}^T \prod_{-m \leq j \leq m; j \neq 0} P(w_{t+j} | w_t; \theta)$$

The objective function $J(\theta)$ is the (average) negative log likelihood:

$$J(\theta) = -\frac{1}{T} \log L(\theta)$$

$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \log P(W_{t+j} | W_t; \theta)$$

Calculation of $P(W_{t+j} | W_t; \theta)$:

We will use two vectors per word w :

- v_w when w is a center word
- u_w when w is a context word

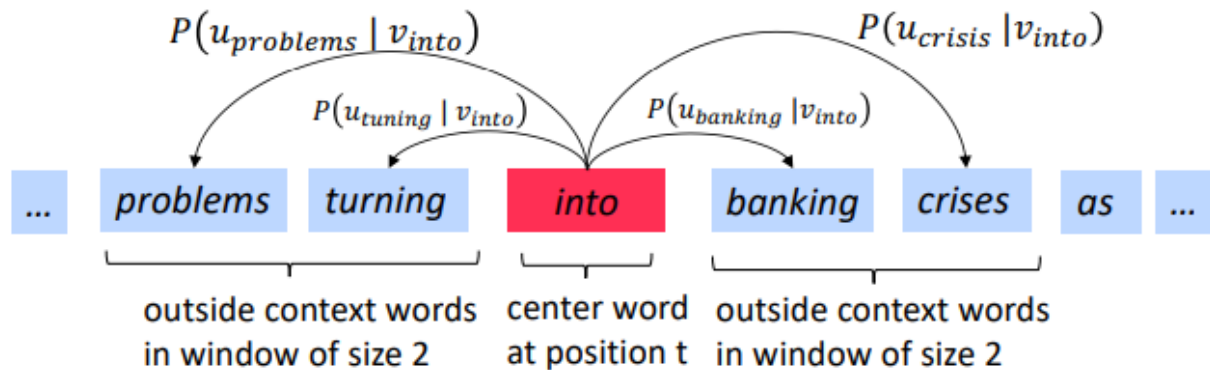
Then for a center word c and a context word o :

$$\text{We know that, } \text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_{j=1}^n \exp(x_j)}$$

$$P(o|c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$

Example:

$P(u_{problems} | v_{into})$ short for $P(problems | into; u_{problems}, v_{into}, \theta)$



Now, θ represents all model parameters, in one long vector. In our case with d -dimensional vectors and V -many words:

$$\theta = \begin{bmatrix} v_{aardvark} \\ v_a \\ \vdots \\ v_{zebra} \\ u_{aardvark} \\ u_a \\ \vdots \\ u_{zebra} \end{bmatrix} \in \mathbb{R}^{2dV}$$

So every word has two vectors. We optimize these parameters by walking down the gradient.

For gradient descent we will have to do calculation of partial derivative with respect to u_o, v_c . We are going to calculate the partial derivative of cost function $J(\theta)$ with respect to v_c :

$$\begin{aligned}\frac{\delta}{\delta v_c} P(o|c) &= \frac{\delta}{\delta v_c} \log \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)} \\ \Rightarrow \frac{\delta}{\delta v_c} \log \exp(u_o^T v_c) - \frac{\delta}{\delta v_c} \log \sum_{w \in V} \exp(u_w^T v_c) \\ \Rightarrow \frac{\delta}{\delta v_c} (u_o^T v_c) - \frac{\delta}{\delta v_c} \log \sum_{w \in V} \exp(u_w^T v_c)\end{aligned}$$

Now, for the calculation of $\frac{\delta}{\delta v_c} (u_o^T v_c)$ we might consider to partially derive with respect to single component of the vector, v_c .

$$\begin{aligned}\frac{\delta}{\delta v_{c_1}} (u_{o_1} v_{c_1} + u_{o_2} v_{c_2} + \dots + u_{o_{100}} v_{c_{100}}) \\ \Rightarrow u_{o_1}\end{aligned}$$

$$\text{So, } \frac{\delta}{\delta v_c} (u_o^T v_c) = u_o$$

Now,

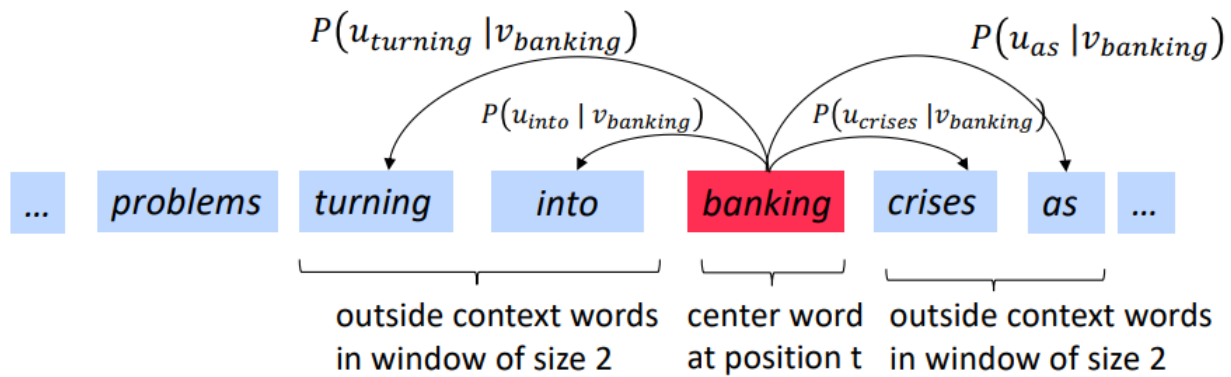
$$\begin{aligned}\frac{\delta}{\delta v_c} \log \sum_{w \in V} \exp(u_w^T v_c) &= \frac{1}{\sum_{w \in V} \exp(u_w^T v_c)} \frac{\delta}{\delta v_c} \sum_{w \in V} \exp(u_w^T v_c) \\ &= \frac{1}{\sum_{w \in V} \exp(u_w^T v_c)} \sum_{w \in V} \frac{\delta}{\delta v_c} \exp(u_w^T v_c) \\ &= \frac{1}{\sum_{w \in V} \exp(u_w^T v_c)} \sum_{w \in V} \exp(u_w^T v_c) \frac{\delta}{\delta v_c} (u_w^T v_c) \\ &= \frac{1}{\sum_{w \in V} \exp(u_w^T v_c)} \sum_{w \in V} \exp(u_w^T v_c) u_w\end{aligned}$$

Hence,

$$\begin{aligned}\frac{\delta}{\delta v_c} P(o|c) &= u_o - \frac{\sum_{w \in V} \exp(u_w^T v_c) u_w}{\sum_{w \in V} \exp(u_w^T v_c)} \\ \frac{\delta}{\delta v_c} P(o|c) &= u_o - \sum_{w \in V} \frac{\exp(u_w^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)} u_w \\ \frac{\delta}{\delta v_c} P(o|c) &= u_o - \sum_{w \in V} P(o|c) u_w\end{aligned}$$

We went through gradient for each center vector v in a window, similarly we can find gradient for each outside vectors u .

Generally in each window we will compute updates for all parameters that are being used in that window.



$$\theta = \begin{bmatrix} v_{\text{aardvark}} \\ v_a \\ \vdots \\ v_{\text{zebra}} \\ u_{\text{aardvark}} \\ u_a \\ \vdots \\ u_{\text{zebra}} \end{bmatrix} \in \mathbb{R}^{2dV}$$

Finally two vector of the same word is averaged to get final word to vector values- $\frac{v_{\text{aardvark}} + u_{\text{aardvark}}}{2}$


There are two model variant:

1. **Skip-grams (SG)** Predict context ("outside") words (position independent) given center word. Example: Dataset would be
(quick, the),
(quick, brown),
(brown, quick),
(brown, fox),
2. **Continuous Bag of Words (CBOW)**: Predict center word from (bag of) context words. Example: Dataset would be
([the, brown], quick),
([quick, fox], brown),
([brown, jumped], fox),

There are two training methods:

1. **Negative sampling** defines an objective by sampling negative examples.
2. **Hierarchical softmax** defines an objective using an efficient tree structure to compute probabilities for all the vocabulary.

Skip Gram Network:



Algorithm Steps:

1. We generate our one hot input vector $x \in \mathbb{R}^{|V|}$ of the center word.
2. We get our embedded word vector for the center word $v_c = Vx \in \mathbb{R}^n$
3. Generate a score vector $z = Uv_c$.
4. Turn the score vector into probabilities, $\hat{y} = \text{softmax}(z)$. Note that $\hat{y}_{c-m}, \dots, \hat{y}_{c-1}, \hat{y}_{c+1}, \dots, \hat{y}_{c+m}$ are the probabilities of observing each context word.
5. We desire our probability vector generated to match the true probabilities which is $y^{(c-m)}, \dots, y^{(c-1)}, y^{(c+1)}, \dots, y^{(c+m)}$ the one hot vectors of the actual output.